

Quiz 2

Nadezhda Gesheva
18 Feb 2018

Text problems

1. Explain in your words what the `unnest_token` function does.

The `unnest_token` function explodes or splits each item from a selected column to one-item-per column frame.

2. Explain in your words what the `gutenbergr` package does.

The `gutenbergr` package provides public access to various texts from the Project Gutenberg (works no longer patented). It also contains metadata for each work, author and subject.

3. Explain in your words how sentiment lexicon work.

Sentiment lexicon provides us with the functionality to obtain the overall sentiment of a text. We do that by comparing the words most frequently used to words which we regards positive/negative (“bing” lexicon); by employing various words expressing emotions such as joy, sadness, fear, etc. (“nrc” lexicon); or by scaling them between -5 and 5 showing overall positiveness/neutrality/negativity of the words in a text.

4. How does `inner_join` provide sentiment analysis functionality.

Inner join enables us to match the words that both exist in a certain lexicon and the words present in the text.

5. Explain in your words what `tf-idf` does.

Tf stands for term frequency and represents the frequency in which words occur in a text. Idf stands for inverse documented frequency, which is a useful technique that gives you actual value that you get out of a token within a document. It does that by assigning weights to the words that are used frequently everywhere, while boosting the weight of frequent words used only in particular chunks.

6. Explain why you may want to do tokenization by bigram.

Bigrams are a useful tool to detect the bias that you might get from text analysis using monograms. For instance, the word “happy” might be the most frequent within a text, and hence we would conclude that the text is positive. Nevertheless, the author of this particular text might use the bigram “not happy” instead of “unhappy”. Hence, our initial conclusion might be wrong.

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##   filter, lag
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.2
```

```
## v ggplot2 3.1.0      v readr  1.1.1
## v tibble  1.4.2      v purrr  0.2.5
## v tidyr   0.8.1      v stringr 1.3.1
## v ggplot2 3.1.0      v forcats 0.3.0

## -- Conflicts ----- tidyverse_conflicts
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(tidytext)
library(gutenbergr)
library(ggplot2)
```

Most frequent words

```
wilde_works <- gutenberg_download(c(174, 301, 773, 774, 790, 844, 854, 873, 875, 885, 887,
                                   902, 921, 1017, 1031, 1057, 1141, 1308, 1338, 14062,
                                   14240, 14522, 23229, 26494, 30120, 30191, 33979, 41806,
                                   42104, 42704, 51563),
                                   meta_fields = "title"
                                   )

## Determining mirror for Project Gutenberg from http://www.gutenberg.org/robot/harvest
## Using mirror http://aleph.gutenberg.org

london_works <- gutenberg_download(c( 215, 310, 318, 710, 746, 788, 910, 911, 1029, 1056, 1074,
                                     1075, 1089, 1096, 1160, 1161, 1162, 1163, 1164, 1187, 1208,
                                     1449, 1596, 1655, 1669, 1688, 1730, 2377, 2415, 2416, 2429,
                                     2512, 2545, 4953, 6455, 10736, 11051, 12336, 14449, 14654,
                                     14658, 16257, 18062, 19678, 21936, 21970, 21971, 22104, 48474),
                                     meta_fields = "title"
                                     )

## Warning in .f(.x[[i]], ...): Could not download a book at http://
## aleph.gutenberg.org/1/9/6/7/19678/19678.zip

wilde_works %>%
  count(title)

## # A tibble: 31 x 2
##   title                                     n
##   <chr>                                <int>
## 1 A Critic in Pall Mall: Being Extracts from Reviews and Miscellan~ 6643
## 2 A Florentine Tragedy; La Sainte Courtisane                1636
## 3 A House of Pomegranates                                    3329
## 4 A Woman of No Importance                                    3274
## 5 "A Woman of No Importance\nAudio performance"              124
## 6 An Ideal Husband                                           4464
## 7 Charmides, and Other Poems                                 2064
## 8 Children in Prison and Other Cruelties of Prison Life       403
## 9 De Profundis                                               1500
## 10 Essays and Lectures                                       5062
## # ... with 21 more rows
```

```
# Our tibble contains 21 works of Wilde
```

```
london_works %>%  
  count(title)
```

```
## # A tibble: 48 x 2  
##   title                                     n  
##   <chr>                                <int>  
## 1 A Daughter of the Snows              10360  
## 2 A Son Of The Sun                     6963  
## 3 Adventure                           8007  
## 4 Before Adam                         3711  
## 5 "Brown Wolf and Other Jack London Stories\nChosen and Edited By ~ 6546  
## 6 Burning Daylight                    11769  
## 7 Children of the Frost                5573  
## 8 Dutch Courage and Other Stories     3730  
## 9 Jerry of the Islands                 7131  
## 10 John Barleycorn                     6236  
## # ... with 38 more rows
```

```
# Our tibble contains 48 works of London
```

```
# Wilde - Transform the tibble to a tidy-text dataset
```

```
tidy_wilde <- wilde_works %>%  
  group_by(title) %>%  
  mutate(line_number = row_number()) %>%  
  unnest_tokens(word, text) %>%  
  ungroup()
```

```
# The most frequently used words by Wilde
```

```
tidy_wilde %>%  
  count(word, sort = TRUE) %>%  
  head(10)
```

```
## # A tibble: 10 x 2  
##   word      n  
##   <chr> <int>  
## 1 the    58693  
## 2 of     36992  
## 3 and    32858  
## 4 to     23067  
## 5 a      20222  
## 6 in     16567  
## 7 is     16359  
## 8 that   13411  
## 9 i      12415  
## 10 it    11480
```

```
# mostly stop words
```

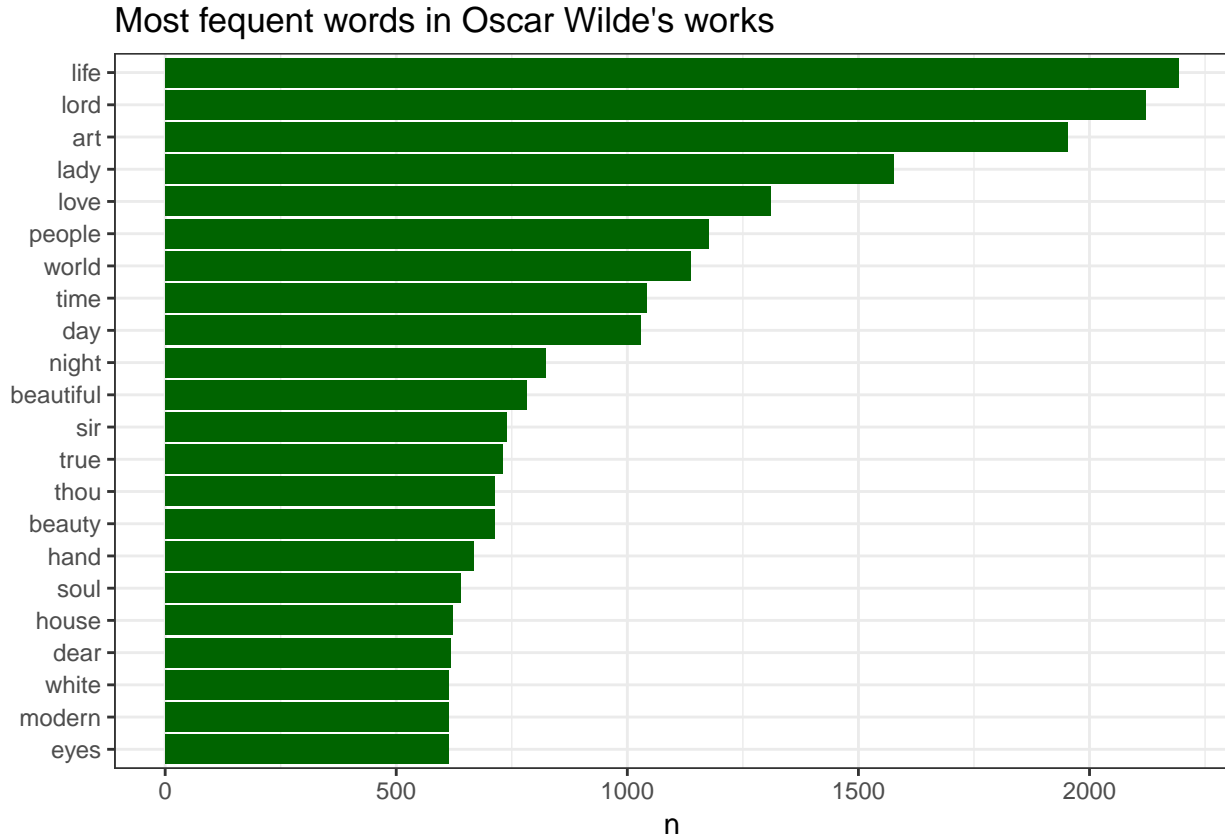
```
tidy_wilde %>%  
  anti_join(stop_words) %>%  
  count(word, sort = TRUE) %>%  
  head(10)
```

```
## Joining, by = "word"
```

```
## # A tibble: 10 x 2
##   word      n
##   <chr> <int>
## 1 life    2193
## 2 lord    2122
## 3 art     1953
## 4 lady    1577
## 5 love    1311
## 6 people  1176
## 7 world   1138
## 8 time    1042
## 9 day     1028
## 10 night   823
```

```
tidy_wilde %>%
  anti_join(stop_words) %>%
  count(word, sort = TRUE) %>%
  filter(n > 600) %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(word, n)) +
  geom_col(fill = "darkgreen") +
  xlab(NULL) +
  ggtitle("Most frequent words in Oscar Wilde's works") +
  theme_bw() +
  coord_flip()
```

```
## Joining, by = "word"
```



```
# London - Transform the tibble to a tidy-text dataset
tidy_london <- london_works %>%
  group_by(title) %>%
  mutate(line_number = row_number()) %>%
  unnest_tokens(word, text) %>%
  ungroup()
```

```
# The most frequently used words by London
tidy_london %>%
  count(word, sort = TRUE) %>%
  head(10)
```

```
## # A tibble: 10 x 2
##   word      n
##   <chr> <int>
## 1 the    184905
## 2 and    113750
## 3 of     80128
## 4 to     65104
## 5 a      61256
## 6 he     48596
## 7 was    47827
## 8 in     45733
## 9 i      41912
## 10 it    35354
```

```
# mostly stop words
```

```
tidy_london %>%
  anti_join(stop_words) %>%
  count(word, sort = TRUE) %>%
  head(10)
```

```
## Joining, by = "word"
```

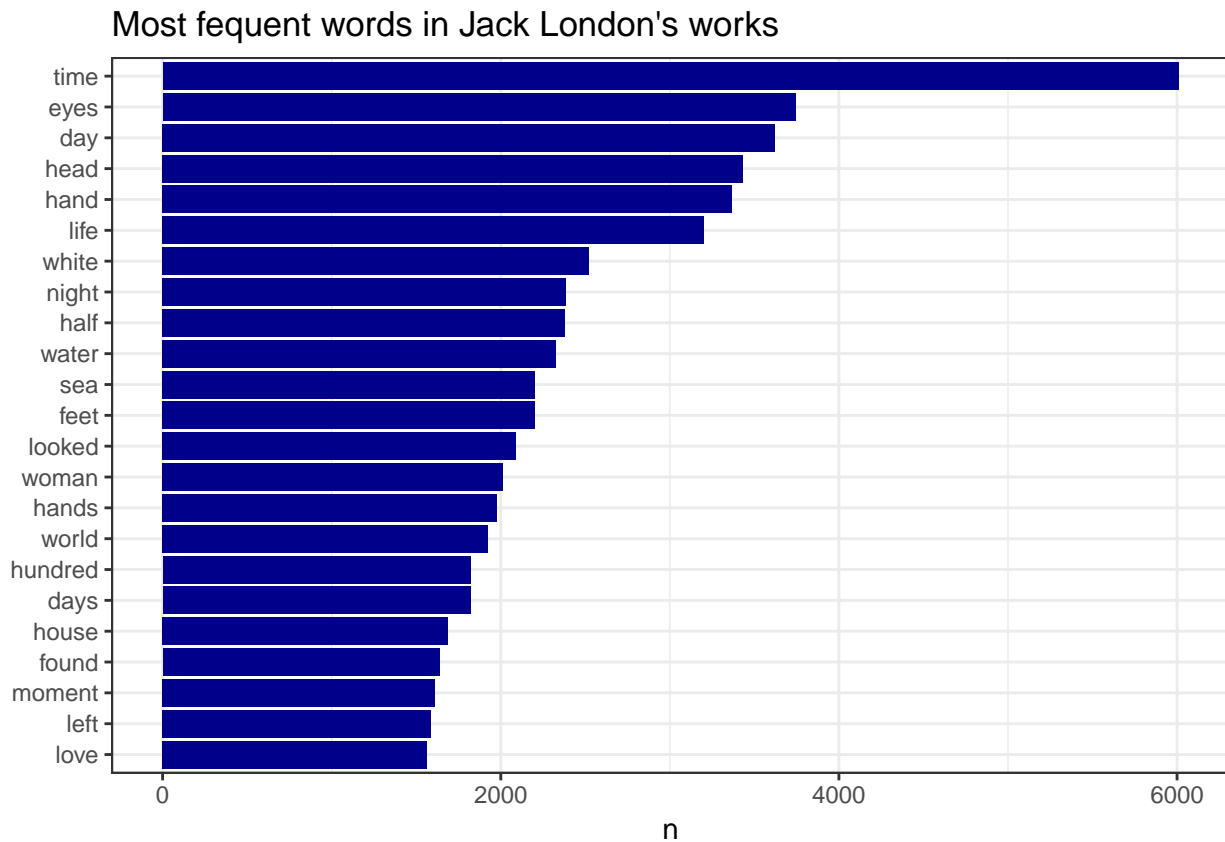
```
## # A tibble: 10 x 2
##   word      n
##   <chr> <int>
## 1 time    6009
## 2 eyes    3743
## 3 day     3624
## 4 head    3432
## 5 hand    3368
## 6 life    3203
## 7 white   2521
## 8 night   2386
## 9 half    2381
## 10 water  2326
```

```
# Life, time and night are in both of top ten most frequent words used by both Oscar and Jack.
```

```
tidy_london %>%
  anti_join(stop_words) %>%
  count(word, sort = TRUE) %>%
  filter(n > 1500) %>%
  mutate(word = reorder(word, n)) %>%
```

```
ggplot(aes(word, n)) +
  geom_col(fill = "darkblue") +
  xlab(NULL) +
  ggtitle("Most fequent words in Jack London's works") +
  theme_bw() +
  coord_flip()
```

```
## Joining, by = "word"
```



Sentiment Analysis

```
t_wilde <- wilde_works %>%
  unnest_tokens(word, text) %>%
  count(title, word, sort = TRUE) %>%
  ungroup()
```

```
sentim_wilde <- t_wilde %>%
  inner_join(get_sentiments("nrc"))
```

```
## Joining, by = "word"
```

```
fear_words_wilde <- sentim_wilde %>%
  filter(sentiment == "fear") %>%
  group_by(word) %>%
```

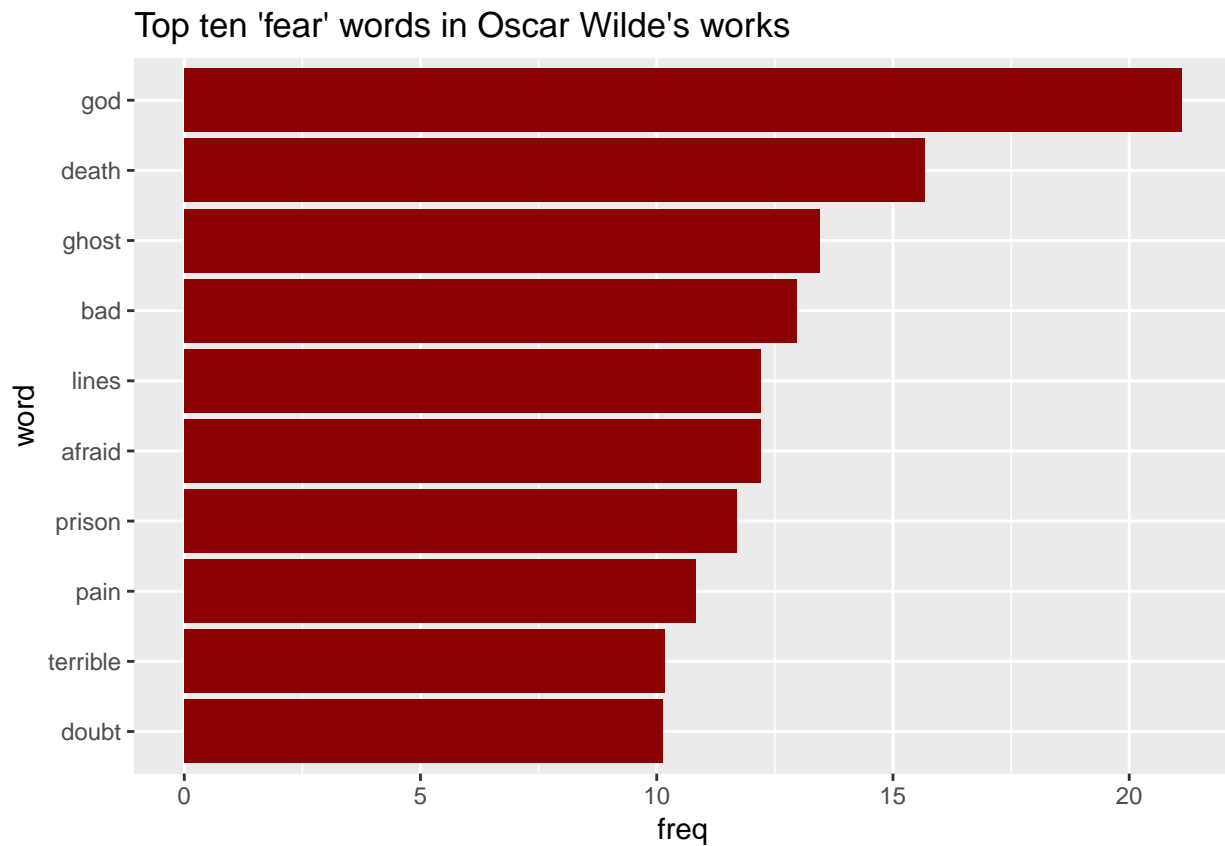
```

summarize(freq = mean(n)) %>%
  arrange(desc(freq))

fear_words_wilde %>%
  top_n(10) %>%
  mutate(word = reorder(word, freq)) %>%
  ggplot(aes(word, freq)) +
  geom_col(fill = "darkred") +
  ggtitle("Top ten 'fear' words in Oscar Wilde's works") +
  coord_flip()

```

Selecting by freq



```

# Similarly for Jack London
t_london <- london_works %>%
  unnest_tokens(word, text) %>%
  count(title, word, sort = TRUE) %>%
  ungroup()

sentim_london <- t_london %>%
  inner_join(get_sentiments("nrc"))

```

Joining, by = "word"

```

fear_words_london <- sentim_london %>%
  filter(sentiment == "fear") %>%
  group_by(word) %>%

```

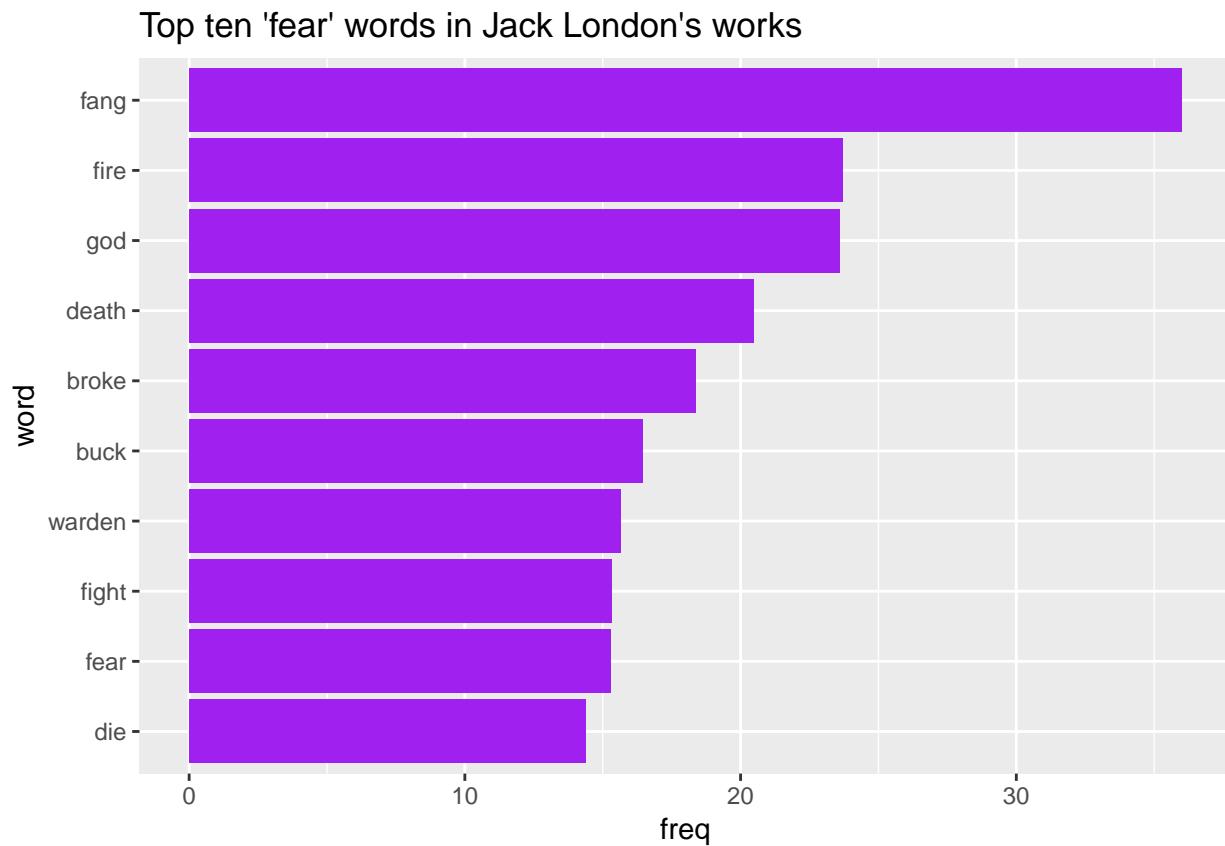
```

summarize(freq = mean(n)) %>%
  arrange(desc(freq))

fear_words_london %>%
  top_n(10) %>%
  mutate(word = reorder(word, freq)) %>%
  ggplot(aes(word, freq)) +
  geom_col(fill = "purple") +
  ggtitle("Top ten 'fear' words in Jack London's works") +
  coord_flip()

```

Selecting by freq



Tf-idf

```

# For the tf-idf analysis we will just pick 10 works per author

# Start with Oscar Wilde
wilde_10_works <- gutenbergl_download(c(174, 301, 773, 774, 790, 844, 854, 873, 875, 885),
  meta_fields = "title")

book_words_wilde <- wilde_10_works %>%
  unnest_tokens(word, text) %>%
  count(title, word, sort = TRUE) %>%

```



```

ungroup()
total_words_wilde <- book_words_wilde %>%
  group_by(title) %>%
  summarize(total = sum(n))

book_words_wilde <- left_join(book_words_wilde, total_words_wilde)

```

```
## Joining, by = "title"
```

```
#book_words_wilde
```

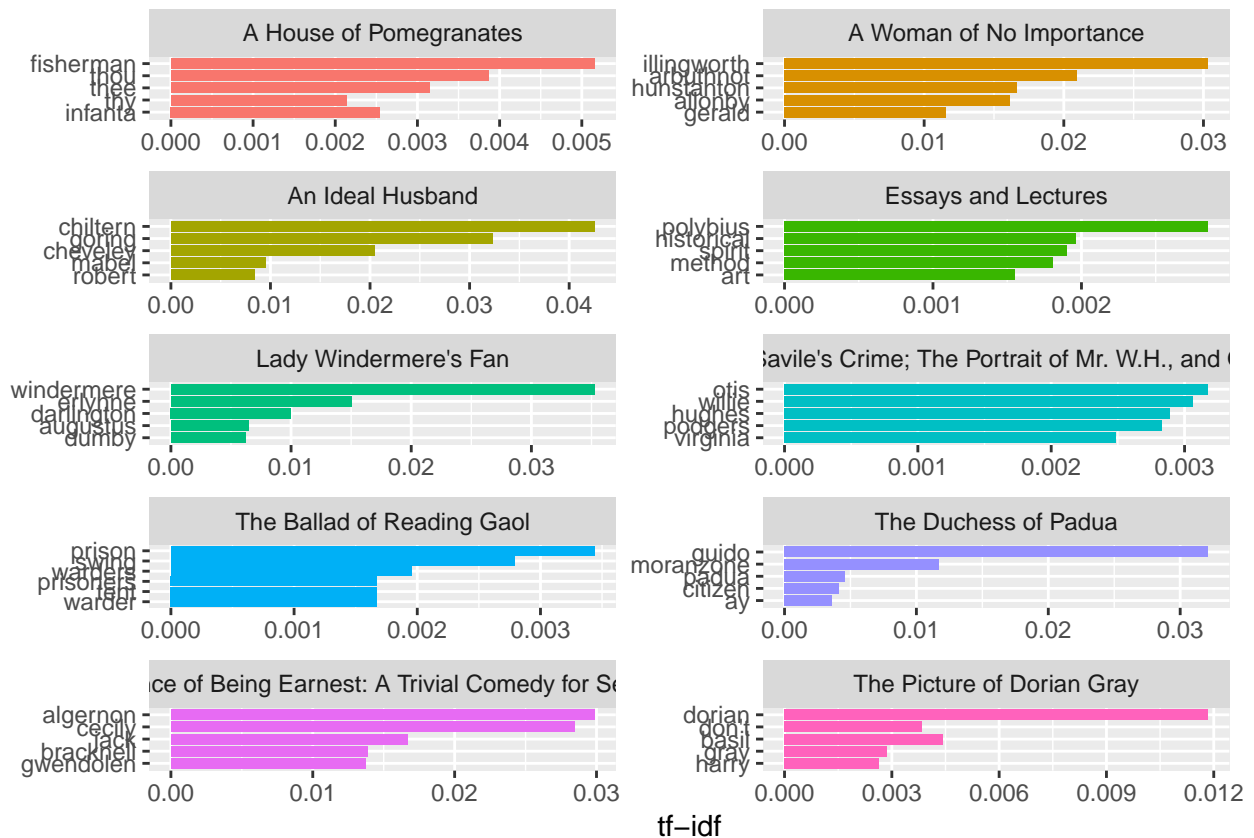
```

book_words_wilde <- book_words_wilde %>%
  bind_tf_idf(word, title, n)

book_words_wilde %>%
  arrange(desc(tf_idf)) %>%
  mutate(word = factor(word, levels = rev(unique(word)))) %>%
  group_by(title) %>%
  top_n(5) %>%
  ungroup %>%
  ggplot(aes(word, tf_idf, fill = title)) +
  geom_col(show.legend = FALSE) +
  labs(x = NULL, y = "tf-idf") +
  facet_wrap(~title, ncol = 2, scales = "free") +
  coord_flip()

```

```
## Selecting by tf_idf
```



```

# Similar analysis tf-idf for Jack London's works

london_10_works <- gutenbergs_download(c( 215, 310, 318, 710, 746, 788, 910, 911, 1029, 1056),
                                       meta_fields = "title")

book_words_london <- london_10_works %>%
  unnest_tokens(word, text) %>%
  count(title, word, sort = TRUE) %>%
  ungroup()
total_words_london <- book_words_london %>%
  group_by(title) %>%
  summarize(total = sum(n))

book_words_london <- left_join(book_words_london, total_words_london)

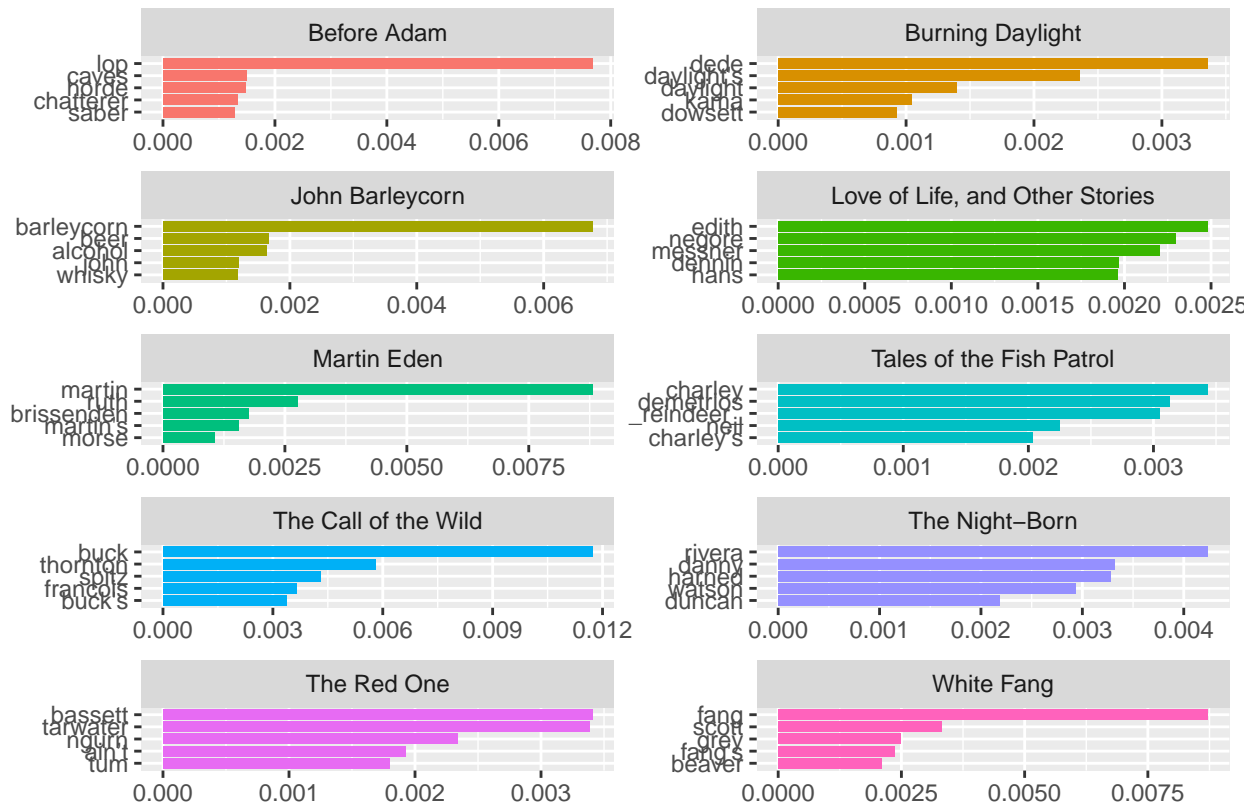
## Joining, by = "title"
#book_words_london

book_words_london <- book_words_london %>%
  bind_tf_idf(word, title, n)

book_words_london %>%
  arrange(desc(tf_idf)) %>%
  mutate(word = factor(word, levels = rev(unique(word)))) %>%
  group_by(title) %>%
  top_n(5) %>%
  ungroup %>%
  ggplot(aes(word, tf_idf, fill = title)) +
  geom_col(show.legend = FALSE) +
  labs(x = NULL, y = "tf-idf") +
  facet_wrap(~title, ncol = 2, scales = "free") +
  coord_flip()

## Selecting by tf_idf

```



tf-idf